

# Introduction to Shared Memory Computing using OpenMP

Pascal Paschos Ph.D.  
Sr. HPC Specialist,  
Research Computing Services


**Northwestern**  
INFORMATION TECHNOLOGY

---

## In this workshop


1. We will learn about elements shared memory architecture and programming
2. Introduce the OpenMP API and it's core components
3. We will use the C-language as the tool to look into code snippets and examples

## The presentation aims


1. To initiate novice programmers to threaded parallelism
  2. To explain the initial means to parallelize serial code
  3. Compile, build & run an openmp executable
- 

# Shared Memory Architecture

- Symmetric multi-processor (SMP) -uniform memory access
  - Shared memory address space with equal access by all CPUs
- NUMA (non-uniform memory access) multi-processor
  - Shared memory address space but time to access can vary by location -Near & Far memory



**Uniform Memory Access**



**Non-Uniform Memory Access**

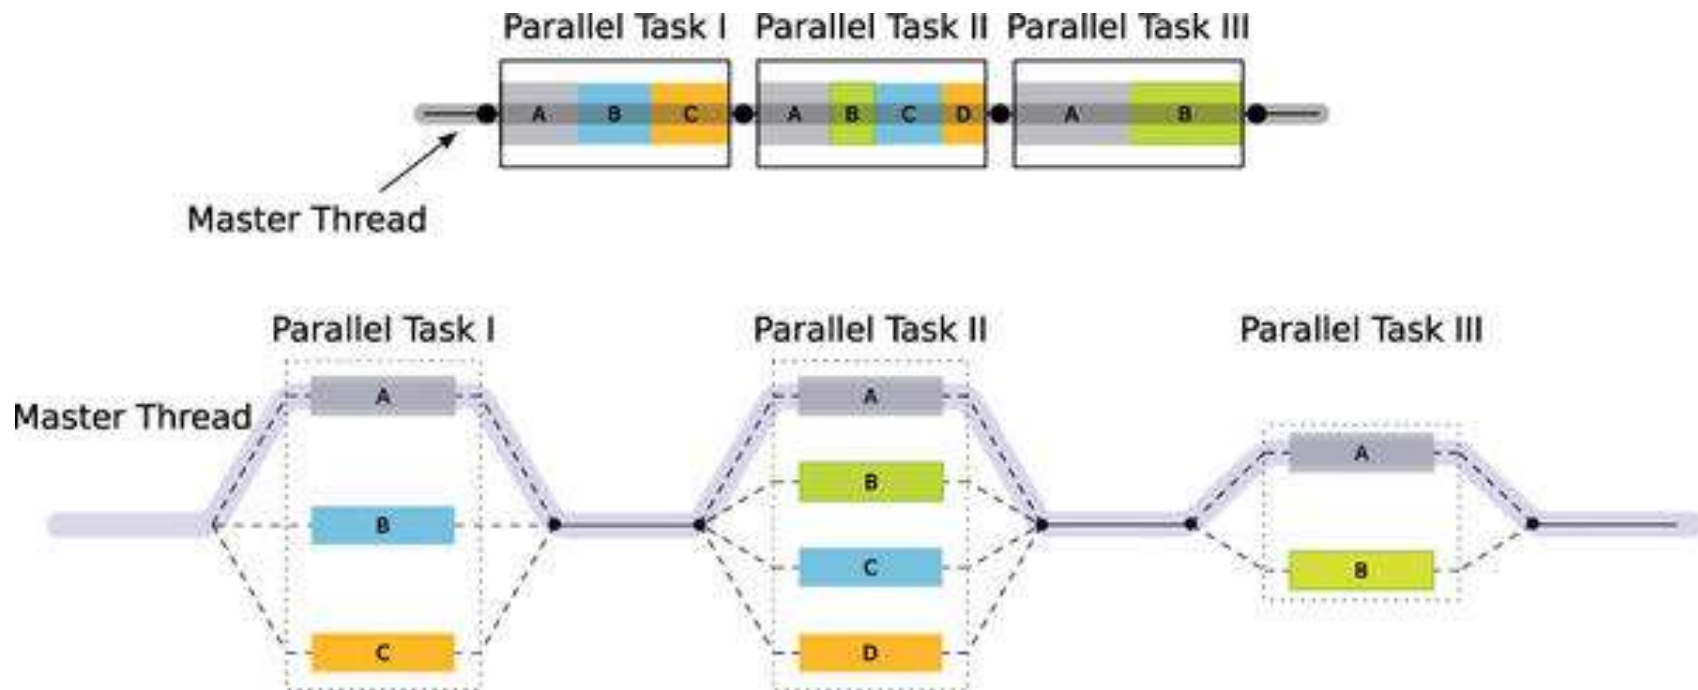
# Shared Memory Programming

- An instance of a program execution is a single main (heavy) process
- Threads are lightweight processes that are spawned by the main process
- All threads share the state of the main process
- Threads 'communicate'/interact via reads & writes to the shared memory address space

# OpenMP Overview

- OpenMP is an API that supports multithreaded, shared memory parallelization
- Simple to use
- Minimally invasive to parallelize serial program
- Cross platform
- Supports Fortran, C/C++

# Fork/join parallelism



# Invoking OpenMP extensions at compilation

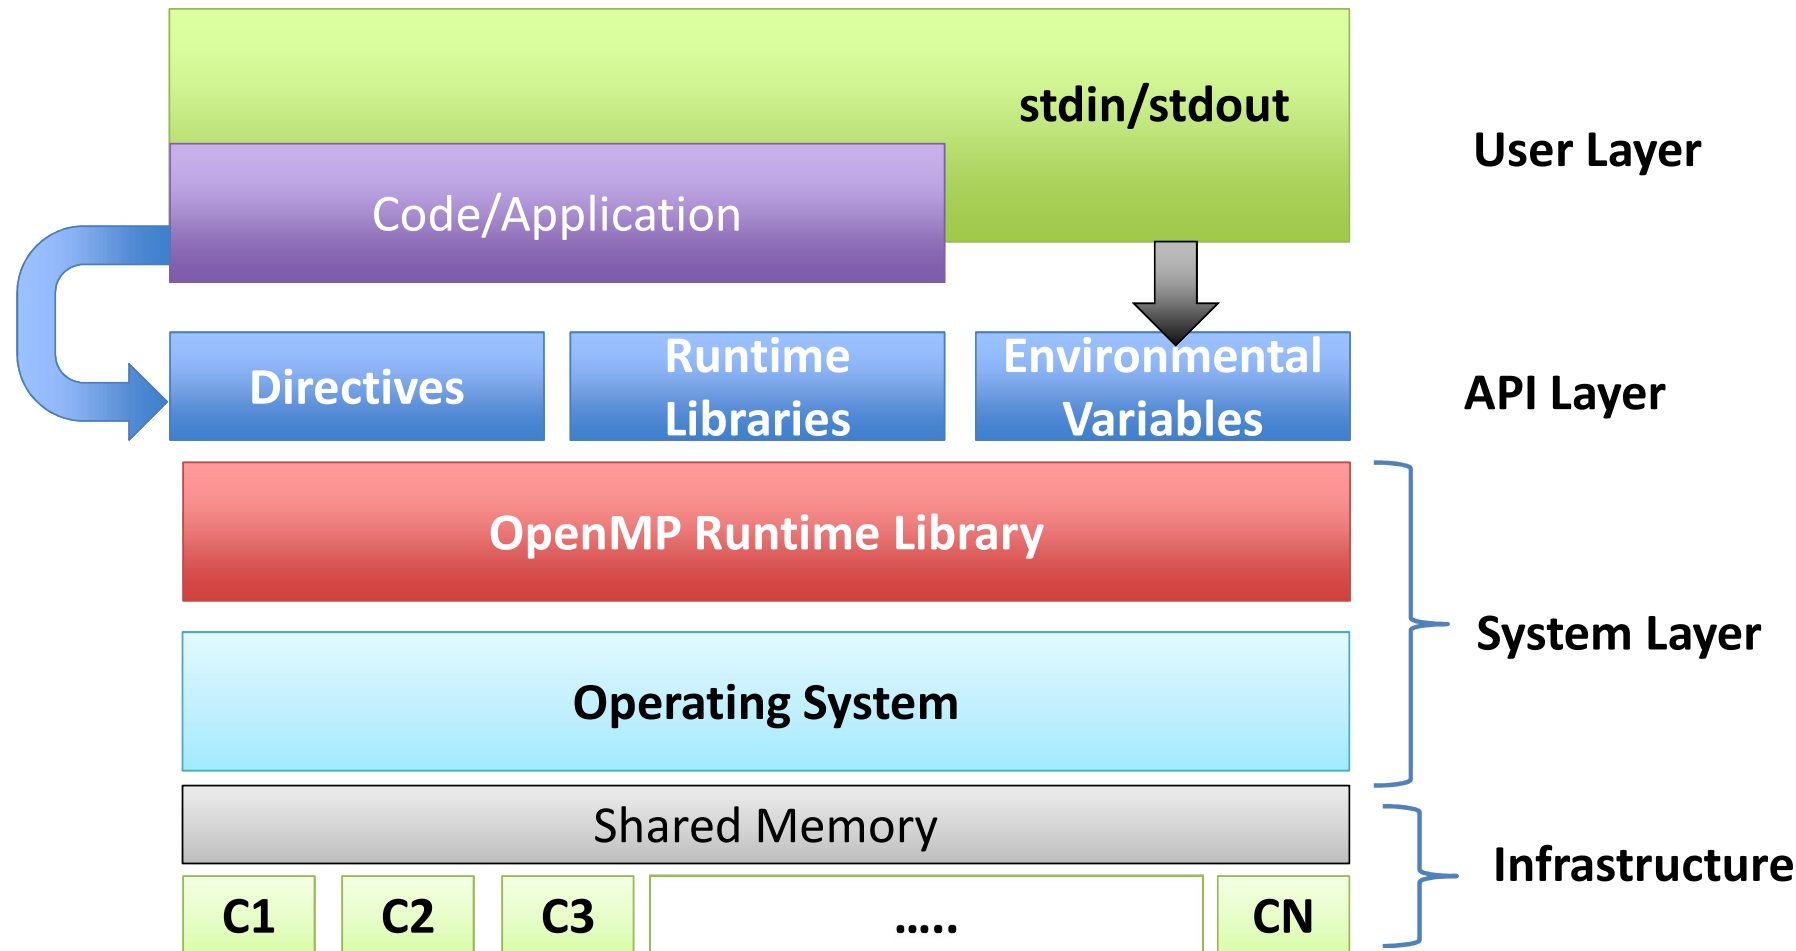
- GNU: `gcc -fopenmp <sourcecode>`
- Intel: `icc -openmp <sourcecode>`
- Must `#include <omp.h>` header file in C/C++ codes
- In a Makefile typically insert in:
- `CFLAGS="$CFLAGS -fopenmp"`



# Components of OpenMP API

- Compiler Directives
- Runtime Library Routines
- Environmental Variables

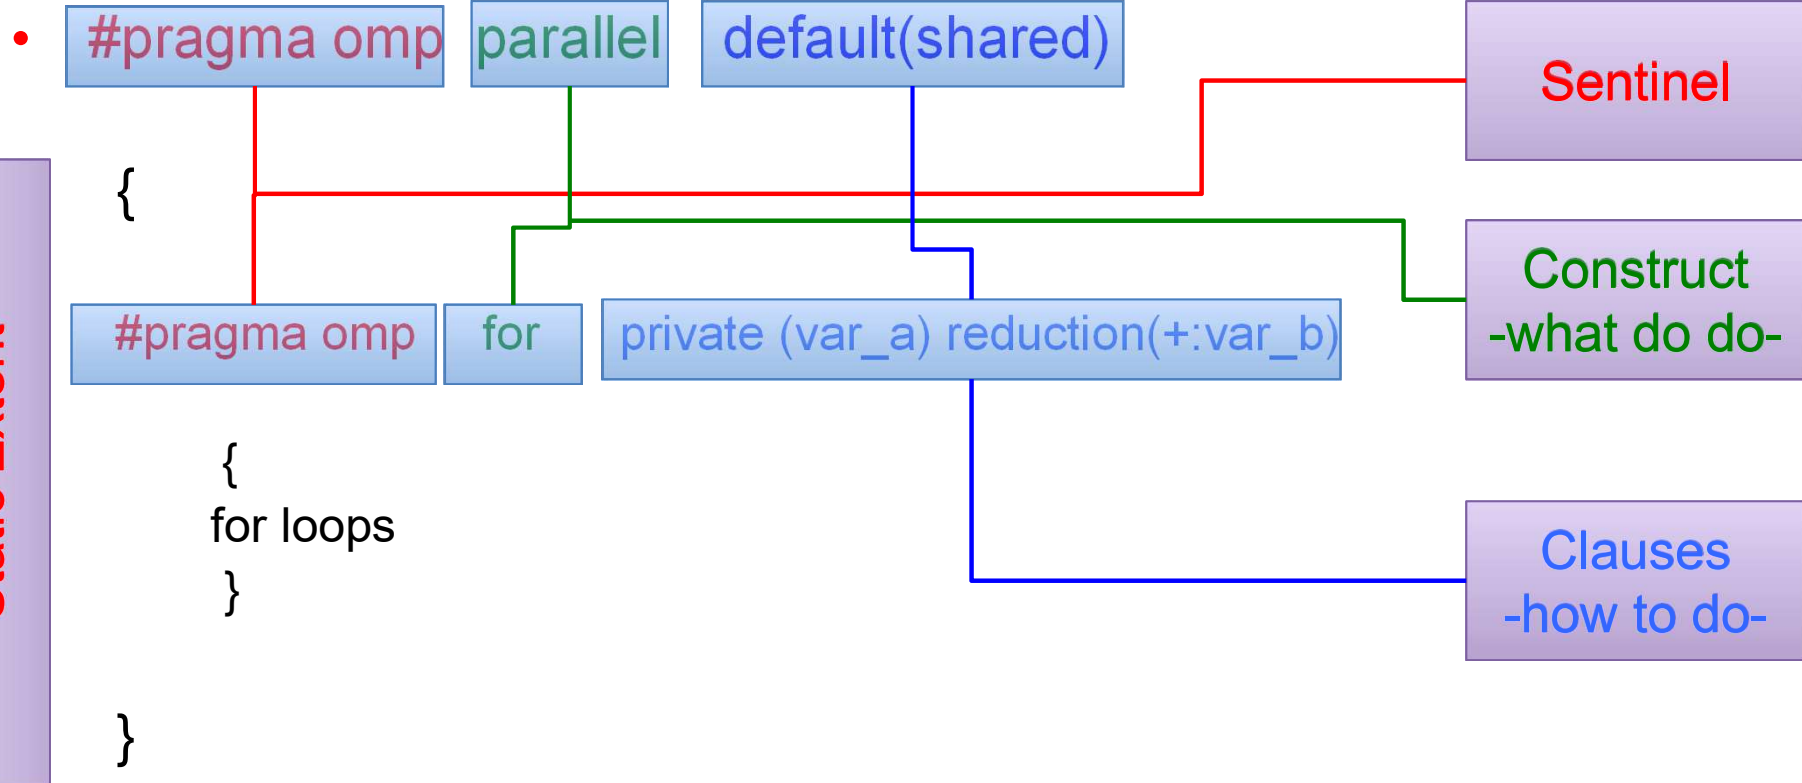
# OpenMP computing stack



# Compiler Directives

- Define parallel regions
- Use constructs for:
  - Work distribution (loops, blocks of code) to threads
  - Synchronization of work among threads
  - Serialization when imperative
- Use clauses to provide tuning and customization

# Generic Compiler Directive Structure

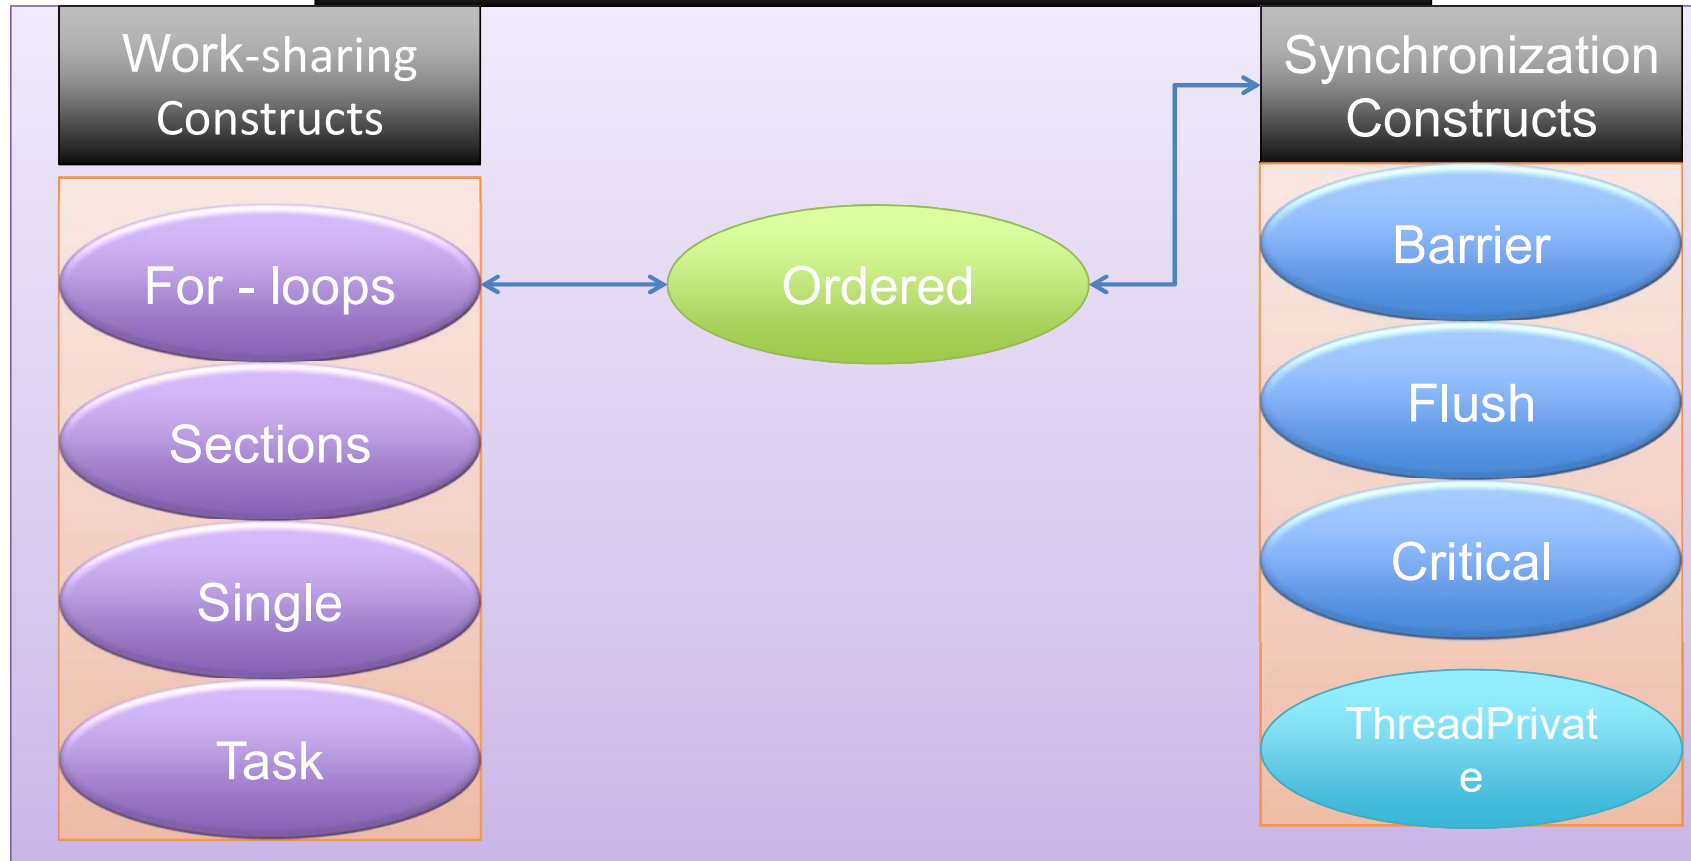


# Defining a parallel region

#pragma omp parallel	<Clauses>
{	if (scalar_expr)
...	private (list)
...	shared (list)
...	default (shared none)
Threaded code here	firstprivate (list)
...	reduction(operator:list)
...	copyin (list)
...	num_threads (int)
}	
End of parallel region	

- All variables defined outside parallel regions are shared by default
- private variables in parallel regions
  - Each thread has it's own copy of the variable and it's own value. The context is hidden from other threads
- shared variables in parallel regions
  - The values of shared variables are visible and accessible by all threads
  - Programmer is responsible for avoiding *race conditions*

## Core Constructs in a Parallel Region C/C++




# for worksharing construct

```
#pragma omp for
```

```
{  
for_loop  
...  
...  
...  
...  
...  
}
```

<Clauses>

```
schedule(type, [chunk])  
ordered  
private (list)  
shared (list)  
firstprivate (list)  
lastprivate (list)  
reduction(operator:list)  
collapse (n)  
nowait
```






The **reduction** clause creates private copies of the shared variable list. Private copies are then ‘reduced’, i.e. summed, and the answer is returned back to the global variable


```
#pragma omp parallel for private(i)  
reduction(+:result)  
{  
    for (i=0; i < n; i++)  
        result = result + i*i;  
}
```

The `schedule(type,chunk)` clause instructs on how to distribute the loop iterations among the threads.

**static:** fixed assignment of loop segments (of size chunk) to threads



**dynamic:** variable assignment of loop segments to threads



The *ordered clause* must be followed by the *ordered construct*. Loop operations are executed sequentially.


```
#pragma omp for ordered  
schedule(dynamic)  
for(int n=0; n<100; ++n)  
{  
    files[n].compress();  
  
    #pragma omp ordered  
    send(files[n]);  
}
```

Done in parallel

Done sequentially

# sections worksharing construct

#pragma omp sections	<Clauses>
<pre>{ #pragma omp section { ... } #pragma omp section { ... } }</pre>	<pre>private (list) shared (list) firstprivate (list) lastprivate (list) reduction(operator:list) <i>nowait</i></pre>



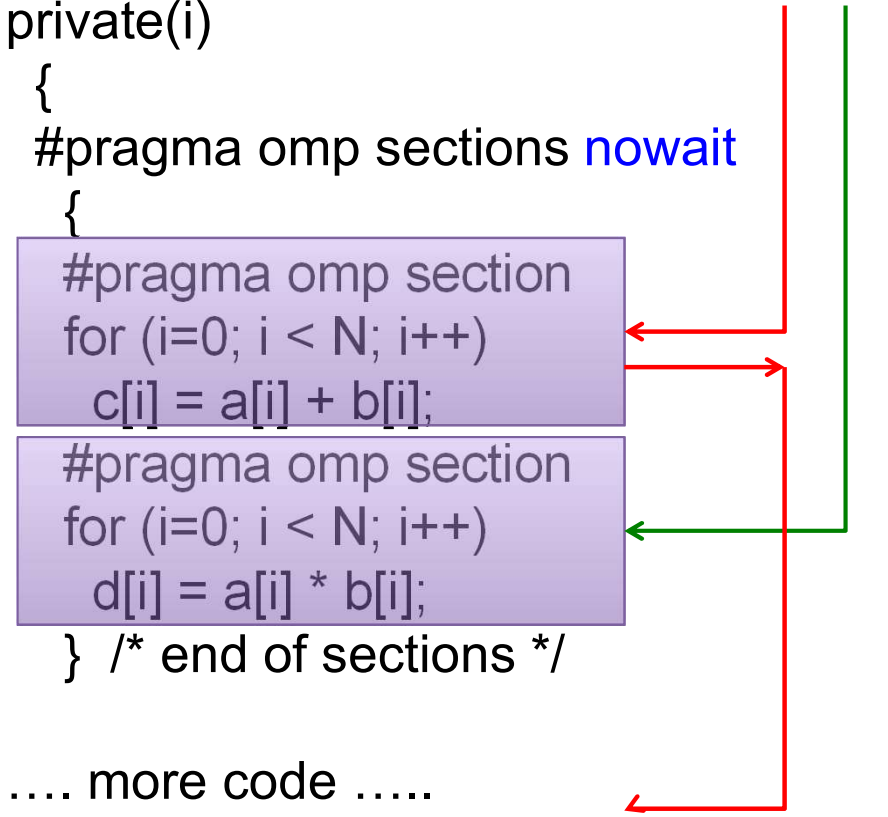
- At the end of 'for', 'sections' & 'single' constructs, threads **by default** are synchronized - they can not proceed unless **all** threads report at the same point in the instruction set.
- The **nowait** clause allows for asynchronous threads.
- Bypassing default synchronization points can lead to incorrect updating of shared(global) variables

```
#pragma omp parallel shared(a,b,c,d)
private(i)
{
    #pragma omp sections nowait
    {
        #pragma omp section
        for (i=0; i < N; i++)
            c[i] = a[i] + b[i];

        #pragma omp section
        for (i=0; i < N; i++)
            d[i] = a[i] * b[i];
    } /* end of sections */


    .... more code .....

} /* end of parallel section */
```



# single “worksharing” construct

#pragma omp single	<Clauses>
{	private (list)
...	firstprivate (list)
block_of_code	nowait
...	
}	



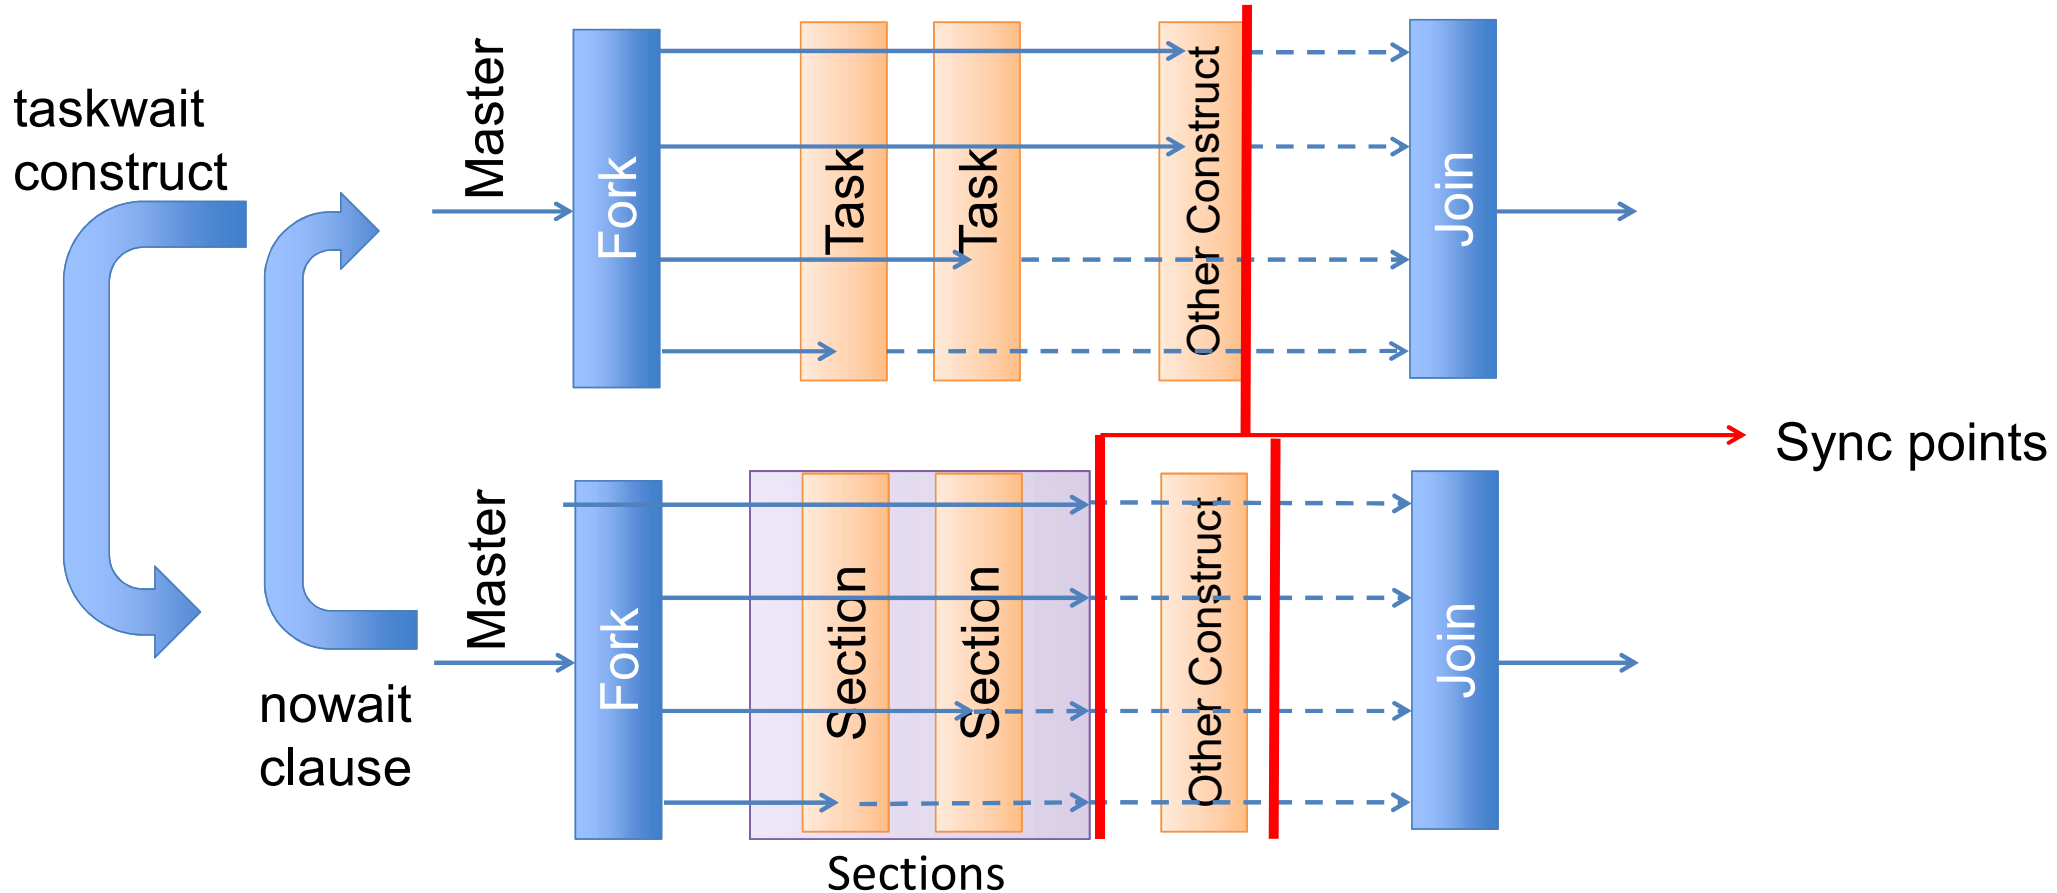
Only a single thread executes code

# Task worksharing construct -OpenMP 3.0

#pragma omp task	<Clauses>
{ block_of_code }	private (list) shared (list) firstprivate (list) if (scalar_expr) untied Default(shared none) final (scalar_expr) mergeable

Synchronization can be forced with `#pragma omp taskwait`

# Task vs. Sections Construct





# Core synchronization directives

```
#pragma omp barrier
```

Barrier synchronizes all threads at that point

```
#pragma omp flush (list)
```

Flush synchronizes all thread-visible (shared) values from local CPU registers to memory. It re-establishes a consistent view of memory across all threads. It is implied upon exiting in:

barrier for  
parallel sections  
critical single  
ordered

```
#pragma omp critical [name]  
structured_block
```

The *critical* directive/construct allows access to a code segment one thread at a time. It aims in protecting simultaneous updates to memory locations of shared variables by threads – before threads are made aware of other updates, i.e. race conditions

```

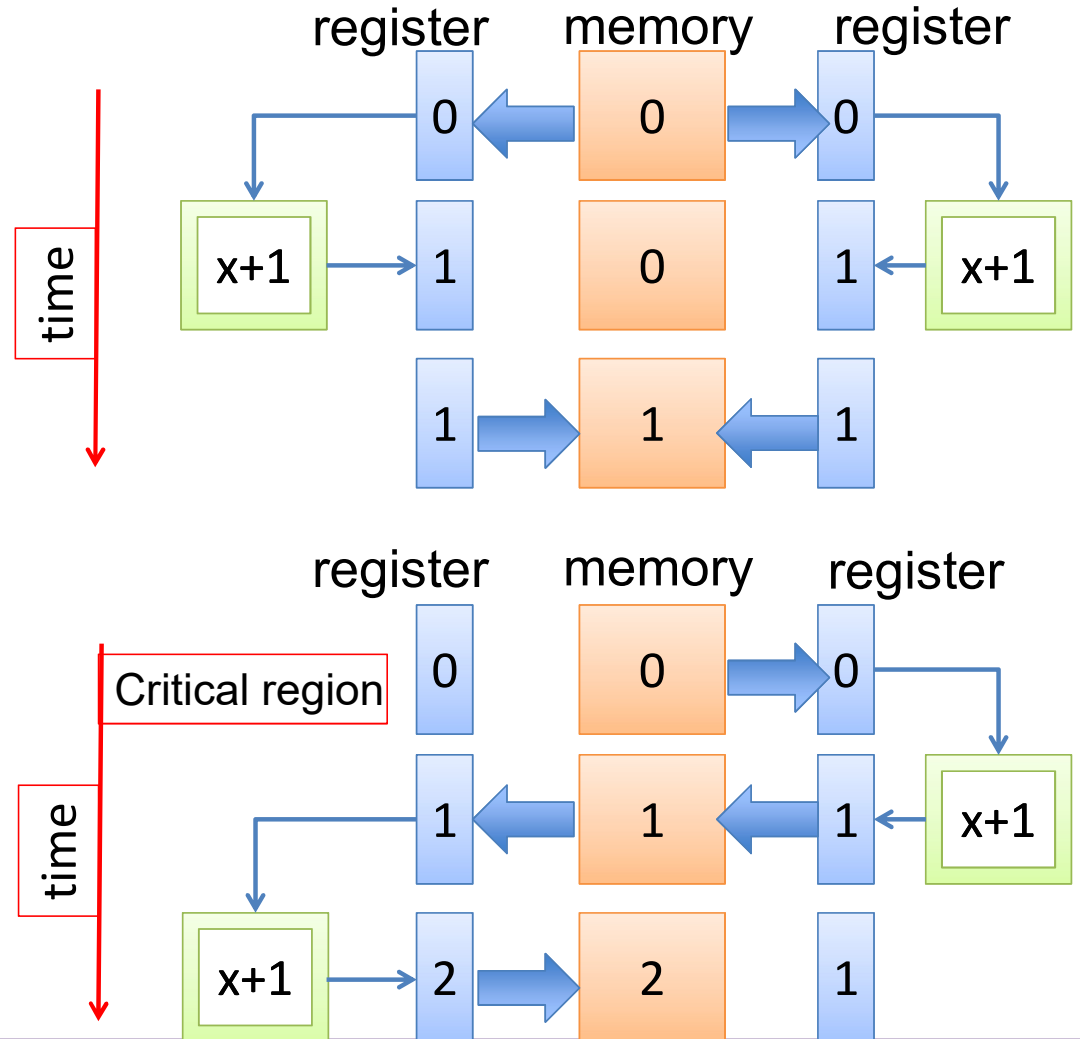
#include <omp.h>
main()
{

int x;
x = 0;

#pragma omp parallel shared(x)
{
    #pragma omp critical
    x = x + 1;
} /* end of parallel section */

}

```



```
#pragma omp threadprivate (list)
```

- The *threadprivate* directive/construct allows private variables to retain their thread specific values across parallel regions
- Does not apply in dynamic schedules
- Number of threads between parallel regions must be the same
- Defined outside the extend of the parallel regions

# Run-time Library Routines

- Explicit management of threads
  - Setting & polling threads
  - Querying thread IDs
  - Querying levels of parallelism (nested threads)
  - Manage thread locking
  - Querying wall clock

# Runtime routines

- Can be executed outside parallel regions
- Routines are hardcoded
- Most frequent examples:
- `tid = omp_get_thread_num()`
  - Query the numeric id number of a thread
- `nthreads = omp_get_num_threads()`
  - Query the number of available threads

# Advantages

Using runtime routines to capture the thread ID allows for:

- Detailed management of thread work and thread synchronization
- Assists in code debugging
- Can emulate practices from distributed computing, i.e. MPI and SPMD (single program multiple data)

# Disadvantages

- Inserts non-pragma source code in the form of runtime routines that prevents code from reverting to pure serial
- Use of SPMD practices requires to significantly alter source code
- Some runtime routines, i.e. `omp_set_num_threads()`, remove code flexibility at execution

# Explicit vs. Implicit Casting to threads

```
int a[nelem],b[nelem];
#pragma omp parallel
{
    int id, i, nthreads, i_begin, i_end;
    id = omp_get_thread_num();
    nthreads = omp_get_num_threads();
    i_begin = id * nelem / nthreads;
    i_end = (id+1) * nelem / nthreads;
    if (id == nthreads-1) i_end = nelems;
    for(i=i_begin;i<i_end;i++) { a[i] = a[i] + b[i];}
}
```

```
int a[nelem],b[nelem];
int i;
#pragma omp parallel for private(i)
for( i=1;i<n_elems;i++)
{a[i] = a[i] + b[i];}
```



# Environment Variables

- Pass information to the executable at runtime, i.e. number of threads or the method of partitioning loop iterations
- Provide an alternative to many runtime routines
- Allows for flexibility at execution time

# Environmental Variable Examples


- Setting the number of threads
  - `export OMP_NUM_THREADS=4`
- Setting the thread stack size
  - `export OMP_STACKSIZE="1G"`
- Binding threads to processors (thread affinity)
  - `export OMP_PROC_BIND=TRUE`
- Specifying how loop operations are partitioned across threads
  - `export OMP_SCHEDULE="static"`

```

#include <stdlib.h>
#include <stdio.h>
#include <omp.h>
#define NSTEPS 500000000
long i,num_steps=NSTEPS;
double x,step,sum,pi;

int main(int argc, char **argv)
{
x=0;
sum = 0.0;
step = 1.0/(double) num_steps;
#pragma omp parallel private(i,x) shared(sum)
{
#pragma omp for schedule(static) reduction(+:sum)
for (i=0; i < num_steps; ++i){
    x = (i+0.5)*step;
    sum = sum + 4.0/(1.0+x*x);
}
}
pi = step * sum;
printf("Computed PI %.24f\n", pi);
return 0;
}

```



$$\pi = \int_0^1 f(x) dx \cong \sum_i f(x_i) \delta x$$

```
[ppk652@quser12 openmp_source]$ gcc -o pi_red.exe -fopenmp pi_red.c
```

```
[ppk652@quser12 openmp_source]$ export OMP_NUM_THREADS=4
```

```
[ppk652@quser12 openmp_source]$ time ./pi_red.exe
```

```
Computed PI 3.141592653589826422688702
```

```
real 0m3.794s
```

```
user 0m14.599s
```

```
sys 0m0.006s
```

```
[ppk652@quser12 openmp_source]$ export OMP_NUM_THREADS=1
```

```
[ppk652@quser12 openmp_source]$ time ./pi_red.exe
```

```
Computed PI 3.141592653590016936959728
```

```
real 0m13.833s
```

```
user 0m13.830s
```

```
sys 0m0.003s
```

```

top - 12:21:06 up 42 days, 23:36, 86 users, load average: 0.84, 0.34, 0.28Tasks: 1128
total, 2 running, 1099 sleeping, 27 stopped, 0 zombie%Cpu(s): 20.2 us, 0.2 sy, 0.0 ni,
79.6 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 stKiB Mem : 13180445+total, 33825440 free,
96267984 used, 1711036 buff/cacheKiB Swap: 10737459+total, 28133764 free,
79240824 used. 33926724 avail Mem  PID USER  PR NI  VIRT  RES  SHR S
%CPU %MEM  TIME+ COMMAND
19861 ppk652  20  0 16796  416  320 R 398.7 0.0 0:14.02 pi_red.exe

```

```

top - 12:22:06 up 42 days, 23:37, 86 users, load average: 0.81, 0.39, 0.30Threads: 3539
total, 5 running, 3259 sleeping, 275 stopped, 0 zombie%Cpu(s): 15.5 us, 0.5 sy, 0.0
ni, 83.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 stKiB Mem : 13180445+total, 33760304 free,
96328160 used, 1715988 buff/cacheKiB Swap: 10737459+total, 28133764 free,
79240824 used. 33864044 avail Mem  PID USER  PR NI  VIRT  RES  SHR S
%CPU %MEM  TIME+ COMMAND
20060 ppk652  20  0 16796  412  320 R 73.1 0.0 0:02.25 pi_red.exe
20061 ppk652  20  0 16796  412  320 R 73.1 0.0 0:02.25 pi_red.exe
20063 ppk652  20  0 16796  412  320 R 73.1 0.0 0:02.25 pi_red.exe
20064 ppk652  20  0 16796  412  320 R 73.1 0.0 0:02.25 pi_red.exe

```

```

#include <stdlib.h>
#include <stdio.h>
#include <omp.h>

int a, b, i, tid;

#pragma omp threadprivate(a)

main () {

/* Explicitly turn off dynamic threads */
omp_set_dynamic(0);

printf("1st Parallel Region:\n");
#pragma omp parallel private(b,tid)
{
tid = omp_get_thread_num();
a = tid;
b = tid;
printf("Thread %d: a,b= %d %d \n",tid,a,b);
} /* end of parallel section */

printf("*****\n");
printf("Master thread doing serial work here\n");
printf("*****\n");

printf("2nd Parallel Region:\n");
#pragma omp parallel private(tid)
{
tid = omp_get_thread_num();
printf("Thread %d: a,b= %d %d \n",tid,a,b);
} /* end of parallel section */

```

```

}

```

## Example of building and running an OpenMP C program with threadprivate directive

```

[ppk652@quser12 openmp_source]$ gcc -fopenmp -o thrpv.exe thrpv.c
[ppk652@quser12 openmp_source]$ export OMP_NUM_THREADS=4
[ppk652@quser12 openmp_source]$ ./thrpv.exe

```

1st Parallel Region:

Thread 2: a,b= 2 2

Thread 3: a,b= 3 3

Thread 0: a,b= 0 0

Thread 1: a,b= 1 1

\*\*\*\*\*

Master thread doing serial work here

\*\*\*\*\*

2nd Parallel Region:

Thread 2: a,b= 2 0

Thread 3: a,b= 3 0

Thread 1: a,b= 1 0

Thread 0: a,b= 0 0

## Instructions for running OpenMP codes on a Mac

From a terminal type: (\$ is the prompt)

```
$/usr/bin/ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
```

```
$brew install gcc --without-multilib
```

Then you can compile from a terminal as: gcc-7 -fopenmp etc

# Conclusions

- The OpenMP API enables thread-based parallelism on shared memory machines using low-level programming languages like C
- Any change in the code requires recompiling
- Parallelism is limited by available memory and number of processors (cores)
- Thread based parallelism is implemented in all programming languages in some form, e.g. Python, R, java



# Threading example in Python

```
import threading

def worker():
    print("Hello world")
    return

threads = []
for i in range(5):
    t = threading.Thread(target=worker)
    threads.append(t)
    t.start()
```

```
[ppk652@quser12 ~]$ python pythread.py
Hello world
Hello world
Hello world
Hello world
```

Questions about OpenMP?

Email us at:

[quest-help@northwestern.edu](mailto:quest-help@northwestern.edu)

Check us out at:

[www.it.northwestern.edu/research](http://www.it.northwestern.edu/research)

